

Kishlay Choudhary

+91-9990687453 | kishlaychoudhary1233@gmail.com | [LinkedIn](#) | [LeetCode](#) | [GitHub](#) | [Portfolio](#)

EDUCATION

Mahatma Gandhi MGM's College of Engineering and Technology
Bachelor in Technology in Computer Science & Engineering

Noida, U.P
Nov. 2021 – June 2025

EXPERIENCE

Research Associate - AI

Keywords Studios

Gurugram, Haryana
August 2025 – August 2026

- Executed 120+ RLHF/HITL evaluation cycles daily (500–600 weekly) on an enterprise foundation LLM using zero-shot, few-shot (3–4 exemplars), and JSON-schema-structured prompting; tracked a reasoning-fidelity metric exceeding a 50% baseline.
- Monitored execution graphs of 2–10 concurrent agents across 100–500+ weekly workflow runs, using state tracking, observation-action loop monitoring, and agent-to-agent communication audits to isolate coordination failures.
- Benchmarked chain-of-thought reasoning and token-efficient JSON-schema tool-call validation, contributing to observed interpretability improvements of approximately 30% and latency reductions approaching 40%, while preventing malformed-output failures.
- Evaluated ETL-style ingestion spanning web extraction, chunk/token transformation, and vector-store loading for an approximately 10M-token context window; validated throughput, persistent cross-session memory, and explicit context-reset overrides.
- Evaluated outputs for harmful, self-harm-related, sexual, and discriminatory content, strengthening safety guardrail coverage and policy compliance across production-bound agentic workflows.
- Authored behavior-analysis reports and reasoning-pattern taxonomies; served as primary India-US liaison for a 5-person research pod across India, Poland, and the US; ranked top-5 for 3 consecutive months and #1 in March 2026.
- Directed RLHF-based behavioral alignment for Mecka AI's humanoid-robotics platform during a 2-month engagement, architecting chained input-output-review workflows across sequential models for industrial-automation and human-companion robot classes.

PROJECTS

AdaptiveRAG | Python, FastAPI, Mistral AI, Qdrant, BM25, Redis, RAGAS

2025

- Architected an adaptive RAG system with Python/FastAPI that classifies queries as simple, complex, or out-of-domain and dynamically routes retrieval; achieved 0.91 RAGAS faithfulness and a 0.86 composite evaluation score on a curated benchmark.
- Engineered hybrid retrieval combining 1024-dimensional Mistral embeddings, Qdrant HNSW vector search, and BM25 through Reciprocal Rank Fusion; added MiniLM cross-encoder reranking, achieving 0.86 context precision and 0.87 answer relevancy.
- Implemented Redis-backed semantic caching with cosine similarity above 0.92 and agentic decomposition for multi-hop queries; used asynchronous FastAPI and SSE streaming to reduce redundant LLM calls and support low-latency responses.

FunnX.AI | Python, Flask, Streamlit, Google Gemini API, OpenRouter, Gunicorn

2024 – 2025

- Engineered a full-stack multi-LLM chat platform with a Flask REST API backend and Streamlit frontend, unifying Google Gemini and DeepSeek within a single workspace.
- Implemented parallel dual-inference “Try Both” mode, dispatching one prompt concurrently through Google Generative AI API and OpenRouter for real-time comparison of reasoning depth, response quality, and speed.
- Built session-aware conversation management preserving context during model switching and a dedicated “Research Mode” prompt configuration for analytical queries; deployed on Render.com with Gunicorn.

TECHNICAL SKILLS

Languages & Backend: Python, Java, FastAPI, Flask, Pydantic, Streamlit, REST APIs, JSON, SSE Streaming

AI & LLM Systems: LangChain, LangGraph, LlamaIndex, OpenAI API, Gemini API, Claude API, Groq, OpenRouter, Mistral AI, Prompt Engineering, Embeddings, RAG Pipelines, Semantic Search, Hybrid Retrieval, BM25, Reciprocal Rank Fusion, Cross-Encoder Reranking, Agentic AI, AI Agents, Tool Calling, Function Calling, Multi-Agent Systems, RLHF, HITL, MCP, RAGAS

Databases & Caching: PostgreSQL (Neon), Redis, Qdrant, FAISS, Pinecone, ChromaDB, Semantic Caching

Developer Tools: Docker, Git, Agile, LLMops, Observability, LangFuse, OpenTelemetry, Gunicorn

Data & Core Concepts: NumPy, Pandas, System Design, DSA, LLM Architecture (Transformers, Attention Mechanism), AI Safety & Governance, AI Alignment

CERTIFICATIONS

Full Stack Generative and Agentic AI with Python

Udemy

Supervised Machine Learning: Regression and Classification

Coursera

FastAPI with GenAI and AgenticAI project - From Basic to AI

Udemy

Python Programming Bootcamp

Udemy